

Cuckoo Cycle: a memory-hard proof-of-work system

John Tromp

January 31, 2014

Abstract

We introduce the first trivially verifiable, scalable, memory-hard and tmt-to-hard proof-of-work system.

1 Introduction

A “proof of work” (PoW) system allows a verifier to check—with negligible effort—that a prover has expended a large amount of computational effort. Originally introduced as a spam fighting measure, where the effort is the price paid by an email sender for demanding the recipient’s attention, they now form one of the cornerstones of crypto-currencies.

Bitcoin[3] uses hashcash[1] as proof of work for new blocks of transactions, which requires finding a nonce value such that twofold application of the cryptographic hash function SHA256 to this nonce (and the rest of the block header) results in a number with many leading 0s. The bitcoin protocol dynamically adjust this “difficulty” number so as to maintain a 10-minute average block interval. Starting out at 32 leading zeroes in 2009, the number has steadily climbed and is currently at 63, representing an incredible $2^{63}/10$ double-hashes per minute. This exponential growth of hashing power is enabled by the highly-parallelizable nature of the hashcash proof of work. which saw desktop cpus out-performed by graphics-cards (GPUs), these in turn by field-programmable gate arrays (FPGAs), and finally by custom designed chips (ASICs).

Downsides of this development include high investment costs, rapid obsolescence, centralization of mining power, and large power consumption. This has led people to look for alternative proofs of work that lack parallelizability, aiming to keep commodity hardware competitive.

Litecoin replaces the SHA256 hash function in hashcash by a single round, 128KB version of the *scrypt* key derivation function. Technically, this is no longer a proof of work system, as verification takes a nontrivial amount of computation. Even so, GPUs are at least an order of magnitude faster than CPUs for Litecoin mining, and ASICs are coming on to the market in 2014.

Primecoin [2] is an interesting design based on finding long Cunningham chains of prime numbers, using a two-step process of filtering candidates by *sieving*, and applying pseudo-primality tests to remaining candidates. The most efficient implementations are still CPU based. A downside to primecoin is that its use of memory is not constrained much.

2 Memory latency; the great equalizer

While cpu-speed and memory bandwidth are highly variable across time and architectures, main memory latencies have remained relatively stable. To level the mining playing field, a proof of work system should be latency bound. Ideally, it should have the following properties:

verify-trivial Implied by the notion of a PoW, but still failed by several systems billed as such.

scalable The amount of memory needed is a parameter that can scale arbitrarily.

linear The number of computational steps and the number of memory accesses is linear in the amount of memory.

tmto-hard There is no time-memory trade-off—using only half as much memory should incur several orders of magnitude slowdown.

random-access RAM is accessed randomly, making bandwidth and caches irrelevant.

parallel-hard Memory accesses cannot be effectively parallelized.

simple The algorithm should be sufficiently simple that one can be convinced of its optimality.

Combined, these properties ensure that a proof of work system is entirely constrained by main memory latency and scales appropriately for any application.

We introduce the first proof of work system satisfying all properties, except for being parallel-resistant at best. Amazingly, it amounts to little more than enumerating nonces and storing them in a hashtable. While all hashtables break down when trying to store more items than it was designed to handle, in one hashtable design in particular this breakdown is of a special nature that can be turned into a concise and easily verified proof. Enter the cuckoo hashtable.

3 Cuckoo hashing

Introduced by Rasmus Pagh and Flemming Friche Rodler in 2001[4], a cuckoo hashtable consists of two same-sized tables each with its own hash function mapping a key to a table location, providing two possible locations for each key. Upon insertion of a new key, if both locations are already occupied by keys, then one is kicked out and inserted in its alternate location, possibly displacing yet another key, repeating the process until either a vacant location is found, or some maximum number of iterations is reached. The latter can only happen once cycles have formed in the *Cuckoo graph*. This is a bipartite graph with a node for each location and an edge for every key, connecting the two locations it can reside at. This naturally suggests a proof of work problem, which we now formally define.

4 The proof of work function

Fix three parameters $L \leq E \leq N$ in the range $\{4, \dots, 2^{32}\}$, which denote the cycle length, number of edges (also easyness, opposite of difficulty), and the number of nodes, respectively. L and N must be even. Function cuckoo maps any binary string h (the header) to a bipartite graph $G = (V_0 \cup V_1, E)$, where V_0 is the set of integers modulo $N_0 = N/2 + 1$, V_1 is the set of integers modulo $N_1 = N/2 - 1$, and E has an edge between $\text{SHA256}(h||n) \bmod N_0$ in V_0 and $\text{SHA256}(h||n) \bmod N_1$ in V_1 for every nonce $0 \leq n < E$. A proof for G is a subset of L nonces whose corresponding edges form an L -cycle in G .


5 Solving a proof of work problem

We enumerate the E nonces, but instead of storing the nonce itself as a key in the Cuckoo hashtable, we store the alternate key location at the key location, and forget about the nonce. We thus maintain the *directed* cuckoo graph, in which the edge for a key is directed from the location where it resides to

its alternate location. The outdegree of every node in this graph is either 0 or 1. When there are no cycles yet, the graph is a *forest*, a disjoint union of trees. In each tree, all edges are directed, directly, or indirectly, to its root, the only node in the tree with outdegree 0. Addition of a new key causes a cycle if and only if its two locations are in the same tree, which we can test by following the path from each location to its root. In case of different roots, we reverse all edges on the shorter of the two paths, and finally create the edge for the new key itself, thereby joining the two trees into one. In case of equal roots, we can compute the length of the resulting cycle as 1 plus the sum of the path-lengths to the node where the two paths first join. If the cycle length is L , then we solved the problem, and recover the proof by enumerating nonces once more and checking which ones formed the cycle. If not, then we keep the graph acyclic by not ignoring the key. There is some probability of overlooking other L -cycles that uses that key, but in the important low easiness case of having few cycles in the cuckoo graph to begin with, it does not significantly affect the rate of solution finding.

6 Implementation and performance

The C-program listed in the Appendix is also available online at <https://github.com/tromp/cuckoo> together with a Makefile, proof verifier and this paper. ‘make test’ tests everything. The main program uses 31 bits per node to represent the directed cuckoo graph, reserving the most significant bit for marking edges on a cycle, to simplify recovery of the proof nonces. On my 3.2GHz Intel Core i5, in case no solution is found, size 2^{20} takes 4MB and 0.25s, size 2^{25} takes 128MB and 10s, and size 2^{30} takes 4GB and 400s, a significant fraction of which is spend pointer chasing.



The left plot above shows the probability of finding a 42-cycle as a function of the percentage edges/nodes, while the right plot shows the average number of memory reads and writes per edge as a function of the percentage nonce/easiness (progress through main loop). Both were determined from 10000 runs at size 2^{20} ; results at size 2^{25} look almost identical. In total the program averages 3.3 reads and 1.75 writes per edge.


7 Memory-hardness

I conjecture that this problem doesn’t allow for a time-memory trade-off. If one were to store only a fraction p of V_0 and V_1 , then one would have to reject a fraction p^2 of generated edges, drastically reducing the odds of finding cycles for $p < 1/\sqrt{2}$ (the reduction being exponential in cycle length). There is one obvious trade-off in the other direction. By doubling the memory used, nonces can be

stored alongside the directed edges, which would save the effort of recovering them in the current slow manner. The speedup falls far short of a factor 2 though, so a better use of that memory would be to run another copy in parallel.

8 Parallelizability

To parallelize the program one could run P processing elements (PE) connected to at least P memory elements (ME) by some multistage interconnection network. For $0 \leq p < P$, PE p processes all nonces $p \bmod P$. Note that unless one uses at least P^2 MEs, memory routing conflicts will frequently arise. Another problem are path conflicts, where one PE is reversing a path that another PE is either following, or installing a new edge on. The left plot below shows the average number of times that either root for a nonce occurred a given number of nonces ago, which suggests that there are potentially about $10P$ such conflicts. Analysing how these conflicts affect performance and the probability of finding cycles is a topic for further research. In any case, development of hardware for improved parallel random access to main memory will benefit other applications too, and so one may expect such improvements to become available on commodity hardware as well.



9 Choice of cycle length

Extremely small cycle lengths risk the feasibility of alternative datastructures that are more memory-efficient. For example, for $L = 2$ the problem reduces to finding a birthday collision, for which a Bloom filter would be very effective, as would Rainbow tables. It seems however that the Cuckoo representation might be optimal even for $L = 4$. Such small values still harm the TMTO resistance though as mentioned in the previous paragraph. In order to keep proof size manageable, the cycle length should not be too large either. We consider 24-64 to be a healthy range, and 42 a nice number close to the middle of that range. The right plot above shows the distribution of cycle lengths found for sizes 2^{10} , 2^{15} , 2^{20} , 2^{25} , as determined from 100000, 100000, 10000, and 10000 runs respectively. The tails of the distributions beyond $L = 100$ are not shown. For reference, the longest cycle found was of length 1726.

10 Scaling memory beyond 16-32 GB

While the current algorithm can accomodate up to $N = 2^{33} - 2$ nodes by a simple change in implementation, a different idea is needed to scale beyond that. To that end, we propose to use K -partite

graphs with edges only between partition k and partition $(k + 1) \bmod K$, where k is fed into the hash function along with the header and nonce. With each partition consisting of at most $2^{31} - 1$ nodes, the most significant bit is then available to distinguish edges to the two neighbouring partitions. The partition sizes should remain relatively prime, e.g. by picking the largest K primes under 2^{31} .

11 Conclusion

Cuckoo Cycle is an elegant memory-hard proof-of-work design featuring scalability, trivial verification, and strong resistance to speedup by special hardware.

References

- [1] Adam Back. Hashcash - a denial of service counter-measure. Technical report, August 2002.
- [2] Sunny King. Primecoin: Cryptocurrency with prime number proof-of-work. July 2013.
- [3] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system. May 2009.
- [4] Rasmus Pagh, Ny Munkegade Bldg, Dk-Arhus C, and Flemming Friche Rodler. Cuckoo hashing, 2001.

12 Appendix A: cuckoo.c Source Code

```
// Cuckoo Cycle, a memory-hard proof-of-work
// Copyright (c) 2013-2014 John Tromp

#include "cuckoo.h"
// algorithm parameters
#define MAXPATHLEN 8192

// used to simplify nonce recovery
#define CYCLE 0x80000000
int cuckoo[1+SIZE]; // global; conveniently initialized to zero

int main(int argc, char **argv) {
    // 6 largest sizes 131 928 529 330 729 132 not implemented
    assert(SIZE < (unsigned)CYCLE);
    char *header = argc >= 2 ? argv[1] : "";
    printf("Looking for %d-cycle on cuckoo%d%d(\"%s\") with %d edges\n",
           PROOFSIZE, SIZEMULT, SIZESHIFT, header, EASINESS);
    int us[MAXPATHLEN], nu, u, vs[MAXPATHLEN], nv, v;
    for (int nonce = 0; nonce < EASINESS; nonce++) {
        sha256edge(header, nonce, us, vs);
        if ((u = cuckoo[*us]) == *vs || (v = cuckoo[*vs]) == *us)
            continue; // ignore duplicate edges
        for (nu = 0; u; u = cuckoo[u]) {
            assert(nu < MAXPATHLEN-1);
            us[++nu] = u;
        }
        for (nv = 0; v; v = cuckoo[v]) {
            assert(nv < MAXPATHLEN-1);
            vs[++nv] = v;
        }
    }
#ifdef SHOW
    for (int j=1; j<=SIZE; j++)

```

```

        if (!cuckoo[j]) printf("%2d:  ",j);
        else             printf("%2d:%02d ",j,cuckoo[j]);
    printf(" %x (%d,%d)\n", nonce,*us,*vs);
#endif
    if (us[nu] == vs[nv]) {
        int min = nu < nv ? nu : nv;
        for (nu -= min, nv -= min; us[nu] != vs[nv]; nu++, nv++) ;
        int len = nu + nv + 1;
        printf("% 4d-cycle found at %d%%\n", len, (int)(nonce*100L/EASINESS));
        if (len != PROOFSIZE)
            continue;
        while (nu--)
            cuckoo[us[nu]] = CYCLE | us[nu+1];
        while (nv--)
            cuckoo[vs[nv+1]] = CYCLE | vs[nv];
        for (cuckoo[*vs] = CYCLE | *us; len ; nonce--) {
            sha256edge(header, nonce, &u, &v);
            int c;
            if (cuckoo[c=u] == (CYCLE|v) || cuckoo[c=v] == (CYCLE|u)) {
                printf("%2d %08x (%d,%d)\n", --len, nonce, u, v);
                cuckoo[c] &= ~CYCLE;
            }
        }
        break;
    }
    if (nu < nv) {
        while (nu--)
            cuckoo[us[nu+1]] = us[nu];
        cuckoo[*us] = *vs;
    } else {
        while (nv--)
            cuckoo[vs[nv+1]] = vs[nv];
        cuckoo[*vs] = *us;
    }
}
return 0;
}

```

13 Appendix B: cuckoo.h Header File

```

// Cuckoo Cycle, a memory-hard proof-of-work
// Copyright (c) 2013-2014 John Tromp

#include <stdio.h>
#include <stdint.h>
#include <string.h>
#include <assert.h>
#include <openssl/sha.h>

// proof-of-work parameters
#ifndef SIZEMULT
#define SIZEMULT 1
#endif
#ifndef SIZESHIFT
#define SIZESHIFT 20
#endif
#ifndef EASINESS
#define EASINESS (SIZE/2)
#endif

```

```

#ifndef PROOFSIZE
#define PROOFSIZE 42
#endif

#define SIZE (SIZEMULT*(1<<SIZESHIFT))
// relatively prime partition sizes
#define PARTU (SIZE/2+1)
#define PARTV (SIZE/2-1)

// generate edge in cuckoo graph from hash(header++nonce)
void sha256edge(char *header, int nonce, int *pu, int *pv) {
    uint32_t hash[8];
    SHA256_CTX sha256;
    SHA256_Init(&sha256);
    SHA256_Update(&sha256, header, strlen(header));
    SHA256_Update(&sha256, &nonce, sizeof(nonce));
    SHA256_Final((unsigned char *)hash, &sha256);
    uint64_t u64 = 0, v64 = 0;
    for (int i = 8; i--;) {
        u64 = ((u64<<32) + hash[i]) % PARTU;
        v64 = ((v64<<32) + hash[i]) % PARTV;
    }
    *pu = 1 + (int)u64;
    *pv = 1 + PARTU + (int)v64;
}

```